

Dataset description:

Data was collected by web crawler from Polish Car trading website - Otomoto.pl

It consists of around 19708 observations (some of them are duplicated) and 18 variables

Objective:

Main objective of this analysis is to build a model which will be able to correctly classify to which Price-group should given car be a part of. Such model can be useful for car buyers or sellers when they are planning to sell or buy a car.

Since variable 'Price' is continuous It was divided into 3 classes 1-highest price 2-medium price and 3-low price car models.

Variables description:

Price-Price of a car

Brand-Brand of a car

Model-Model of a car

Year_produced-year in which car was produced

Mileage-car's mileage

Cylinders_capacity - Car's cylinders capacity

Fuel_type- Diesel/Petrol etc.

HP - Horse Power

transmission - type of transmission

drive_type - type of drive (FWD,AWD,RWD)

Colour-car's colour

Serviced - whether it was serviced in authorized mechanic

New/Used-Describes whether car was used before or is it brand new

```
data_all.columns
[306] ✓ 0.2s
... Index(['Price', 'Brand', 'Model', 'Year_produced', 'mileage',
        'Cylinders_capacity', 'Fuel_type', 'HP', 'transmission', 'drive_type',
        'liters_per_km', 'Type', 'CO2 emission', 'No_of_doors', 'No_of_seats',
        'Colour', 'Serviced', 'New/Used'],
        dtype='object')
```

Feature engineering:

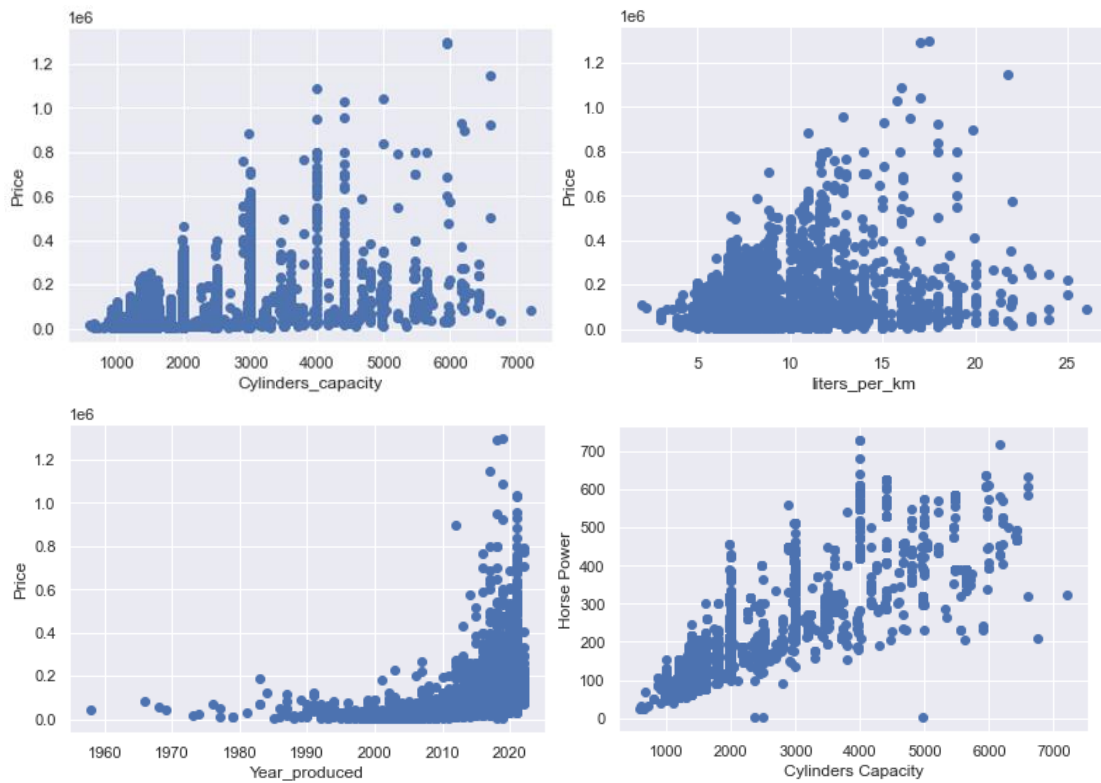
- Removed two outliers (first one because its price was too big and second one because its liters per km ratio was too high)
- Reviewed each variable separately and converted it to an expected type and format
- Filled missing values with mean values of a sub groups
- Removed the rest if there were still missing values
- Grouped car brands and models as 'Other' when they appearing too rarely
- All numeric variables were transformed using log1p transformation due to their skewness

Final 'clean' dataset consists of 8503 observations

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8503 entries, 1 to 19706
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Price                 8503 non-null  float64
1   Brand                 8503 non-null  object
2   Model                 8503 non-null  object
3   Year_produced         8503 non-null  int32
4   mileage               8503 non-null  float64
5   Cylinders_capacity    8503 non-null  float64
6   Fuel_type             8503 non-null  object
7   HP                   8503 non-null  float64
8   transmission          8503 non-null  object
9   drive_type            8503 non-null  object
10  liters_per_km         8503 non-null  float64
11  Type                  8503 non-null  object
12  CO2 emission          8503 non-null  float64
13  No_of_doors           8503 non-null  float64
14  No_of_seats           8503 non-null  float64
15  Colour                8503 non-null  object
16  Serviced              8503 non-null  object
17  New/Used              8503 non-null  object
dtypes: float64(8), int32(1), object(9)
memory usage: 1.2+ MB
```

EDA :

Below few of correlation graphs are presented:



All correlations are as we should expect:

- The higher cylinder capacity the higher the price
- The higher the cylinder capacity the more horse power car has
- The higher the fuel burning rate the higher the price
- The newer the car the more expensive it is

Correlation between variables:

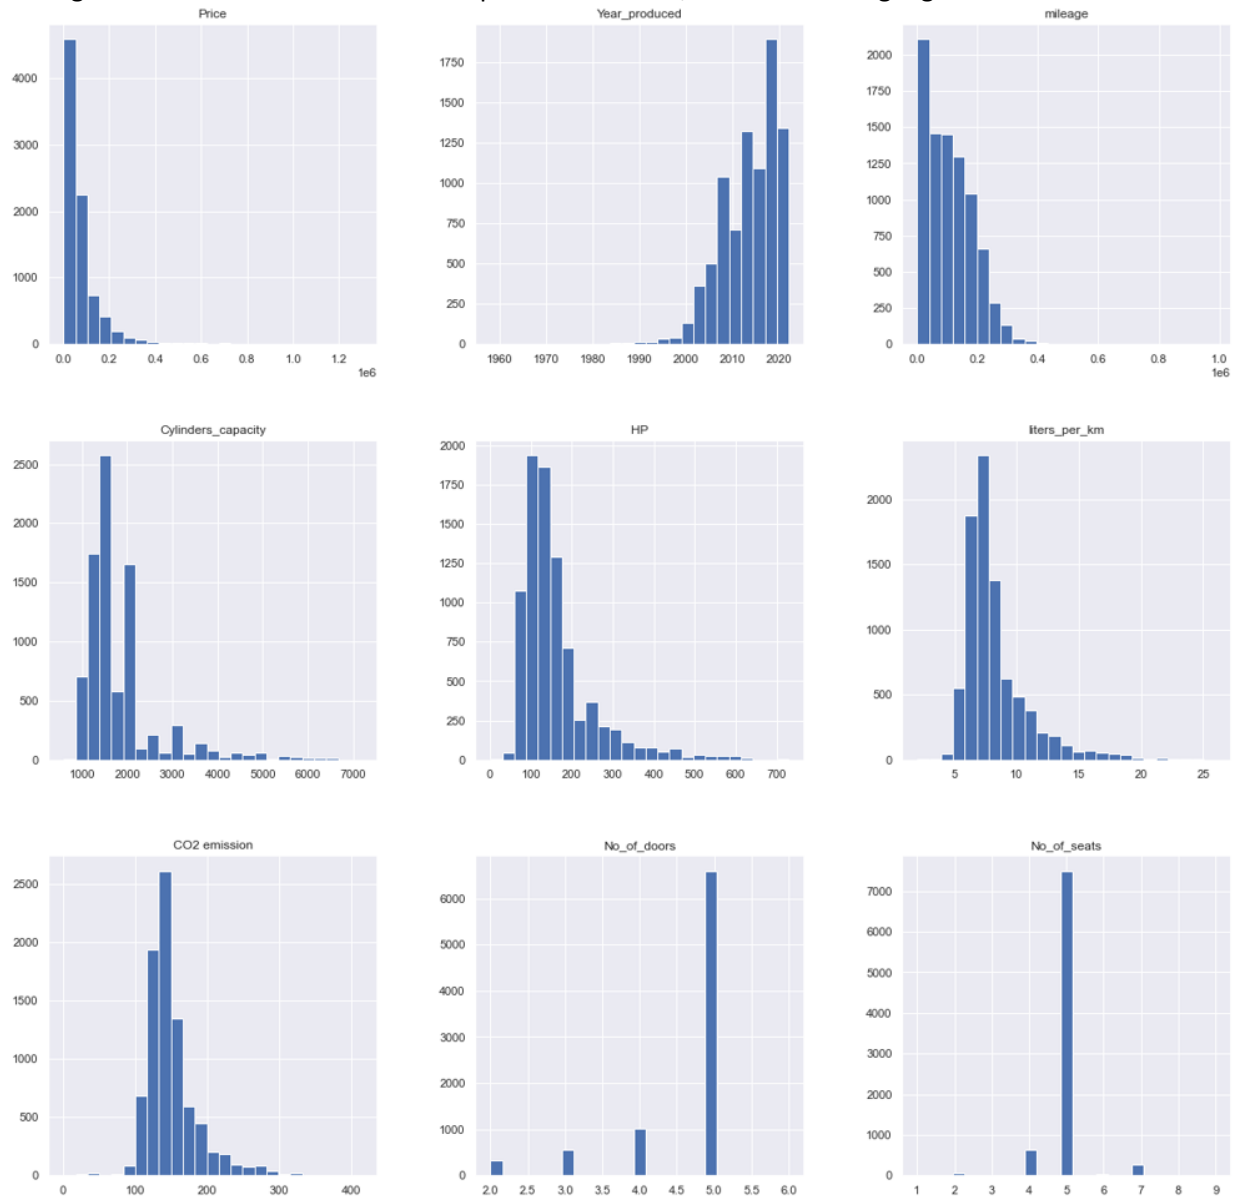
When it comes to direction in which different parameters are correlated with the price:

There are top 30 negatively and positively correlated:


mileage	-0.565052	Price	0.743590
liters_per_km	-0.350799	No_of_doors_5.0	0.220582
transmission_Manual	-0.255964	Colour_White	0.168758
drive_type_RWD	-0.235059	HP	0.150132
CO2_emission	-0.209621	No_of_seats_5.0	0.118961
No_of_doors_3.0	-0.180385	Colour_Gray	0.084075
Fuel_type_Petrol	-0.169636	Model_tipo	0.080530
Cylinders_capacity	-0.168949	Brand_skoda	0.079569
Colour_Silver	-0.134902	Model_q3	0.079486
Model_Other	-0.133237	Brand_dacia	0.075577
Model_sl	-0.105919	Model_arteon	0.072460
Model_a4	-0.096026	Model_xc40	0.071079
Model_seria3	-0.094904	Brand_hyundai	0.063899
No_of_seats_4.0	-0.090991	drive_type_FWD	0.063314
Model_vectra	-0.083228	Model_kuga	0.059911
Colour_Green	-0.073809	Model_stelvio	0.058353
Colour_Other	-0.069774	Model_duster	0.056153
Model_xj	-0.068230	Model_tucson	0.056000
Model_clk	-0.065800	Brand_volvo	0.053735
Model_grandvitara	-0.065063	Brand_kia	0.053553
Model_meriva	-0.063083	Model_spacestar	0.053518
Colour_Gold	-0.059779	Model_xc60	0.053138
No_of_seats_2.0	-0.059733	Serviced_Yes	0.052882
Brand_bmw	-0.059520	Model_superb	0.052706
Model_9-3	-0.059374	Model_compass	0.051220
Brand_saab	-0.059374	Model_tiguan	0.049479
Brand_daihatsu	-0.059018	Model_giulia	0.049156
Brand_honda	-0.058072	Model_ateca	0.048580
Model_klasas	-0.057067	Model_captur	0.047847
Model_accord	-0.055774		
Name: Year_produced, dtype: float64		Name: Year_produced, dtype: float64	

As we can see all of them are more or less in line with common sense


Histograms of continuous variables are presented below, all are indicating high skewness:



Histograms after performing logarithmic transformation look much closer to normal distribution:



Then variable 'Price' was transformed into 3 classes 1- High price 2-Medium price 3-Low price so that classification models can be used to estimate it:



Classification:

Models used:

- Logistic Regression
- Decision Tree Classifier
- Random Forest Classifier
- Gradient Boosting Classifier

Since classes of Target variable varied greatly in size, my primary metric was ROC AUC

models were fit and train on dataset, training and fitting were done using KFold cross validation, on 5 folds, in order to minimize the overfit some parameters have been introduced :

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, roc_auc_score
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import KFold

model_LR= LogisticRegression(solver='liblinear')
model_tree=DecisionTreeClassifier()
model_forest=RandomForestClassifier(n_estimators=500,random_state=42,warm_start=True)
model_xgb=GradientBoostingClassifier(n_estimators=400, learning_rate=0.009)

param_list={}
✓ 0.5s

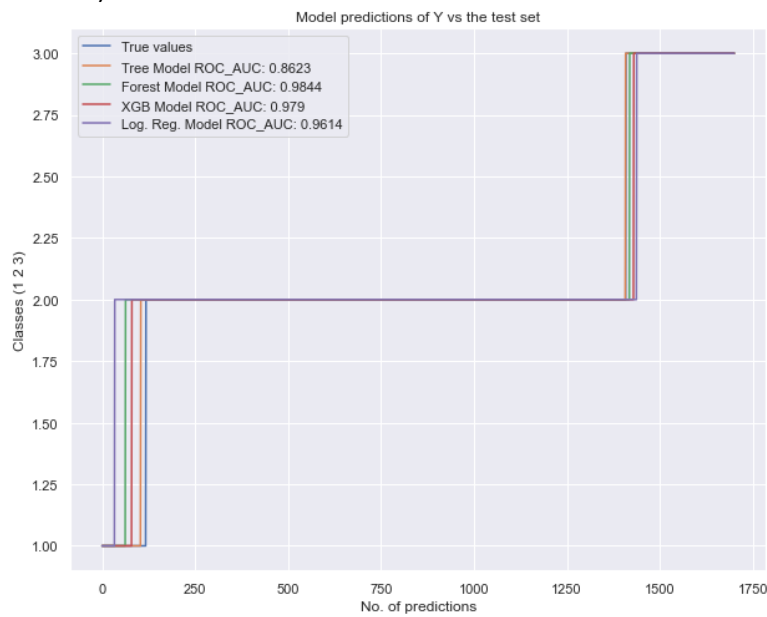
kf=KFold(n_splits=5,shuffle=True,random_state=42)
X_train, X_test, Y_train, Y_test = train_test_split(X_class, Y_class, test_size=0.2,random_state=42,shuffle=True)
```

Results:


▽ Logistic Regression							
					precision	recall	f1-score support
				1	0.74	0.21	0.33 118
				2	0.89	0.97	0.93 1291
				3	0.89	0.80	0.85 292
				accuracy			0.89 1701
				macro avg	0.84	0.66	0.70 1701
				weighted avg	0.88	0.89	0.87 1701
▽ Decision Tree Classifier							
					precision	recall	f1-score support
				1	0.70	0.62	0.66 118
				2	0.94	0.95	0.94 1291
				3	0.87	0.88	0.88 292
				accuracy			0.91 1701
				macro avg	0.84	0.81	0.83 1701
				weighted avg	0.91	0.91	0.91 1701
▽ Random Forest Classifier							
					precision	recall	f1-score support
				1	0.92	0.49	0.64 118
				2	0.93	0.98	0.95 1291
				3	0.92	0.89	0.90 292
				accuracy			0.93 1701
				macro avg	0.92	0.79	0.83 1701
				weighted avg	0.93	0.93	0.92 1701
▽ Gradient Boosting Classifier							
					precision	recall	f1-score support
				1	0.86	0.58	0.70 118
				2	0.93	0.98	0.95 1291
				3	0.92	0.86	0.89 292
				accuracy			0.93 1701
				macro avg	0.91	0.81	0.85 1701
				weighted avg	0.93	0.93	0.92 1701

When it comes to ability of identifying particular classes Random Forest generally yield better results, however f1 score for class 1 was higher in case of XGB model.

Below is the summary of how those models performed on test set (around 1700 values were used as test set):



And test set (around 700 values were used as a train set):



Best model overall in this case is **Random Forest**. However due to its high overfitting on training set (even after initial tries to reduce it) **XGBoost** may present more stable results. However it can be done as a next step on improving the model. To focus on just those two classifying methods and modify the hyperparameters.